

Welcome to Software Carpentry Etherpad!

This pad is synchronized as you type, so that everyone viewing this page sees the same text. This allows you to collaborate seamlessly on documents.

Use of this service is restricted to members of the Software Carpentry and Data Carpentry community; this is not for general purpose use (for that, try [etherpad.wikimedia.org](http://etherpad.wikimedia.org)).

Users are expected to follow our code of conduct: <http://software-carpentry.org/conduct.html>

All content is publicly available under the Creative Commons Attribution License: <https://creativecommons.org/licenses/by/4.0/>

## **Introduction to HPC workshop: 05/31/2017**

**Welcome!**

**Instructor: Andrea Zonca (SDSC)**

**Helpers:**

**Ryan Johnson (Library)**

**Reid Otsuji (Library)**

**Guilherme (Gui) Castelao (SIO)**

**Please sign the sign-in sheet, thank you!**

Sticky notes:

Pink - need help

Yellow - all good

**Setup SSH:**

<https://psteinb.github.io/hpc-in-a-day/setup/>

**Connect to Comet**

Comet training user accounts:

username  
etrainXX (XX number will be assigned by instructor, if you come in late please ask for a number)

password  
tr6nasr2

ssh etrainXX@comet.sdsc.edu (replace XX with your assigned number)

git clone https://github.com/zonca/intro\_hpc

(lesson materials)

work will be done in terminal on Comet

git clone github lesson materials in terminal

pwd  
ls  
cd intro\_hpc

module avail | less  
# lists software packages and versions available to load

q to exit

module list #currently loaded packages

module show python #show details about package

module load python

module load abaqus

# modules temporarily modify your environment  
# modules define paths that you want to use in your terminal during your session, but it's possible to return to this environment later if set up properly

module # lists all available commands  
# To load a specific version, for example:

module avail abaqus # Shows which versions of abaqus are available  
module load abaqus/6.11-2 # Loads abaqus version 6.11-2, which was listed in the previous step as available

# You can install software on your home folder, but the sysadmin installs modules on the supercomputer

cd # make sure you are in your home folder

nano .bashrc

# past command to .bashrc file:

export MODULEPATH=/share/apps/compute/modulefiles/applications:\$MODULEPATH

# save .bashrc

[ctrl] D # to exit comet, then ssh back in again

ssh etrainXX@comet.sdsc.edu (replace XX with your assigned number)

module avail anaconda # Shows that anaconda is now available, and which version it is available.

module load anaconda # Loads anaconda, which will make a newer version of python available

which python # Shows which is the default python.

python # Initialize python and should show which version is running, like "Python 3.6.0 ..."

# good practice to load minimal number of packages

# create a conda environment

conda create --name py --clone root

# failed to create will show, but it's ok - OS error is ok

# Alternative:

conda create -n py --clone=/share/apps/compute/anaconda

# activate and deactivate conda environment

source activate py

conda remove conda -env

conda install line\_profiler

which python

conda list

# Do in this order:

module load anaconda

source activate py

source deactivate

```
conda remove conda-env
conda install line_profiler
```

```
#
cd intro_hpc
cd 1_scheduler
```

```
# view contents of multiple_commands.sh file
cat multiple_commands.sh
```

```
# run multiple_commands.sh file
bash multiple_commands.sh
```

```
# You log into a login node in Comet, then send jobs to the scheduler, which sends jobs to
the computing nodes. The scheduler executes a bash script.
```

```
sbatch multiple_commands.sh # it's ok if you see sbatch error
```

```
# sleep 1 min, verify we are running from scheduler
```

```
sbatch --time=00:03:00 multiple_commands.sh
# minimal necessary to interact with the scheduler, run command multiple_commands for 3
mins - time for jobs to run
# The scheduler will allocate you in the queue based on the time that you requested. Shorter
the time that you request, faster the scheduler can find a slot for you, but if your job goes
longer than that, it will be interrupted without finish.
```

```
# nothing is being executed, only being added to the queue
```

```
# SLURM name of scheduler
```

```
# when you get a job number check squeue , this will show the current jobs
squeue -u $USER
```

```
# Comet user guide
http://www.sdsc.edu/support/user\_guides/comet.html#running
```

```
/home folder available on all nodes
```

```
scancel [job ID number] # kill job e.g. scancel 9486298
```

```
#job .out files will be saved in folder
```

cat slurm-[jobnumber].out # to view job information

# to run HPC job  
create script  
choose time  
run job

#add time to .sh file , create line in .sh file

#SBATCH --time=00:03:00

# check manual for commands  
man scancel  
man squeue

# Let's install some Python libraries that we'll use in the following examples.

source activate py # Getting inside 'py' environment of conda. You know you're  
inside with the line starts with '(py)'.  
conda list # Probably a small list, missing numpy and other important  
libraries. Let's install new packages.  
conda install numpy # Install numpy library in your conda environment.  
# If you get a message error like: "InstallError: Install error: Error: one or more of the  
packages already installed depend on 'conda' ...",  
# you'll need to clean it first with the following command  
conda remove conda-env  
# Now let's stall it for good.  
conda install numpy dask distributed

# conda package manager a good way to install packages

# container technology e.g. Docker  
# Create an OS image on your laptop, load container onto Comet, then run your jobs in your  
container.

Singularity, more at <http://singularity.lbl.gov/quickstart>

# run singularity\_shell.sh  
bash singularity\_shell.sh

cat /etc/\*release

# /home dir available in singularity container

# Exit the singularity container by Ctrl+D

# You know that you're outside the container if your command line starts with the '(py)' instead of 'Singularity....'

cd 1\_scheduler

ls # You should see multiple\_commands\_singularity.sh in this directory

# Now run:

sbatch --time 0:03:00 multiple\_commands\_singularity.sh &

# Note the '&' in the end of the command that put this command line to run background.

squeue -u \$USER # Shows everything that you're running. It should return something close to:

```
Submitted batch job 9487906
      JOBID PARTITION  NAME  USER ST   TIME  NODES NODELIST
(REASON)
      9487906  compute multiple etrain95 PD   0:00    1 (None)
```

# Check inside multiple\_commands\_singularity.sh, the last line is:

singularity exec \$IMAGE bash multiple\_commands.sh

# Note the 'exec' in the previous line. It's saying to run the script multiple\_commands.sh with bash inside the image defined by \$IMAGE, using singularity. Because the exec, it is not opening the virtual machine, but using that image to run multiple\_commands.sh and return.

cd ../2\_tasks/

mkdir files

# The script create\_input\_files.sh will create empty files. Let's run it

bash create\_input\_files.sh

python process\_file.py files/data\_0000.txt # This is a fake process procedure for our class.  
It waits for a short time pretending

# that is doing something

partition=compute: 24 cores -> 128GB

partition=shared: 1 core -> 5MB

--ntasks-per-node=n (n = number of cores)

# view .slrm file

cat 2\_submit\_for.slrn

[http://www.sdsc.edu/support/user\\_guides/comet.html#storage](http://www.sdsc.edu/support/user_guides/comet.html#storage)

```
# view slrm array file
cat 3_submit_slurmarray.slm  # parallelizes the job
```

```
# submit job
sbatch 3_submit_slurmarray.slm
```

```
# view job queue
squeue -u $USER
```

```
# google 'slurm job array' for more information
```

```
https://slurm.schedmd.com/job\_array.html
```

To loop over files, create a list of files and then loop with slurm array

```
#==== AFTERNOON SECTION =====
```

```
Examples to run applications:
  /share/apps/examples
```

```
ibrun to run MPI applications (see AMBER)
```

```
# home file system info
http://www.sdsc.edu/support/user\_guides/comet.html#storage
```

```
# do not store data on home file system
# home file system for storing code, etc., is backed up
```

```
# storing scratch data
```

```
# store data in $USER <-- [your username]
```

```
Lustre Comet scratch filesystem: /oasis/scratch/comet/$USER/temp_project (parallel file
system)
```

```
# To create a soft link, run:
```

```
ln -s /oasis/scratch/comet/$USER/.... name_of_the_alias
# Doing that, you would be able to access the /oasis/... by the shortcut name_of_the_alias
```

```
# to show available accounts
show_accounts
```

```
#SBATCH -A <<project>>
```

SSD Scratch space

/scratch/\$USER/\$SLURM\_JOB\_ID (212 GB for compute, shared partition)

# To visualize the next example:

cd ~/intro\_hpc/3\_filesystems

nano access\_ssd.slrn

www.globus.org

# transfer large files between supercomputers

copying files

scp <file> user@comet.sdsc.edu:/<path>

google: ssh without password

Python help: eg. np.random.uniform? (shows help about method)

# Explanation of pi estimate

# The area of the circle is pi times the square of radius, i.e.:

$$A = \pi * r^2$$

# If we re-arrange the equation

$$A / r^2 = \pi$$

# Which means, the area of the circle (A) divided by the area of the square that contains this circle ( $r * r$ ) is equal to pi

# To simplify, let's do this in one quarter of the image, and in the end multiply the ratio by 4.

Parallel processing in Python, use: import concurrent.futures

# To launch Jupyter notebook:

cd ~/intro\_hpc

sbatch 5\_notebook\_dask/notebook.slrn

# To cancel all jobs:

scancel -u \$USER

Setting a password:

[https://github.com/zonca/intro\\_hpc/blob/master/5\\_notebook\\_dask/create\\_password.shf](https://github.com/zonca/intro_hpc/blob/master/5_notebook_dask/create_password.shf)

# Token is in the .out file



#

launch Jupyter in browser, e.g.,  
[http://comet-03-18.sdsc.edu:8888/tree?](http://comet-03-18.sdsc.edu:8888/tree?token=1ffb9803bf28385b7bb8cd4f0542088df49853bc3ceb20b7)  
[token=1ffb9803bf28385b7bb8cd4f0542088df49853bc3ceb20b7](http://comet-03-18.sdsc.edu:8888/tree?token=1ffb9803bf28385b7bb8cd4f0542088df49853bc3ceb20b7)

# There are wifi connection issues with:  
UCSD-GUEST  
Need to use UCSD-PROTECTED

# Safari and Internet Explorer will not load Py kernel

Can also ssh into compute node:  
e.g., ssh comet-04-19 -> interactive shell

In Jupyter:  
To list files: !ls <shift><enter>

n\_samples\_inside

4 \* n\_samples\_inside / 3000

list(map(inside\_circle, [1000, 1000, 1000]))

# Parallel processing:

```
import concurrent.futures
n_workers = 3
executor = concurrent.futures.ProcessPoolExecutor(
    • max_workers = n_workers)
np.sum(list(map(inside_circle, [1000, 1000, 1000, 1000])))
```

# in jupyter terminal try:  
time python 1\_serial\_digits\_of\_pi.py 100000 12

# notice time difference for parallel processing

Using dask array:

replace:  
x = np.random.uniform(size=total\_count)

with:

```
x = da.random.uniform(size=total_count, chunks=total_count//48)
```

##### Supercomputers at SDSC #####

Comet (international system) & TSCC

Can request TSCC trial

[http://www.sdsc.edu/services/hpc/hpc\\_systems.html](http://www.sdsc.edu/services/hpc/hpc_systems.html)

[xede.org](http://xede.org)

conda install ipykernel

ipython kernel install --user --name=py